

Unified Transactions Reconciliation Platform (UTRP)

SQL Analysis

-- Creating a small warehouse for development

```
CREATE OR REPLACE WAREHOUSE UTRP_WH  
  WAREHOUSE_SIZE = 'XSMALL'  
  AUTO_SUSPEND = 60  
  AUTO_RESUME = TRUE;
```

-- Creating database and schema

```
CREATE OR REPLACE DATABASE UTRP_DB;  
CREATE OR REPLACE SCHEMA UTRP_SCHEMA;
```

-- Creating Tables

```
CREATE OR REPLACE TABLE CUSTOMERS(  
  ext_customer_key VARCHAR,  
  name VARCHAR,  
  email VARCHAR,  
  created_at DATE  
);
```

```
CREATE OR REPLACE TABLE INVOICES(  
  ext_invoice_key VARCHAR,  
  ext_customer_key VARCHAR,  
  invoice_date DATE,  
  amount NUMBER(10, 2),  
  currency VARCHAR,  
  status VARCHAR,  
  source_system VARCHAR,  
  inserted_at TIMESTAMP_NTZ  
);
```

```
CREATE OR REPLACE TABLE PAYMENTS(  
  ext_payment_key VARCHAR,  
  ext_customer_key VARCHAR,  
  payment_date DATE,  
  amount NUMBER(10, 2),  
  currency VARCHAR,  
  payment_method VARCHAR,  
  source_system VARCHAR,  
  inserted_at TIMESTAMP_NTZ  
);
```

-- Data Preview

```
SELECT * FROM CUSTOMERS;  
SELECT * FROM INVOICES;  
SELECT * FROM PAYMENTS;
```

-- Performing validation of the data

```
SELECT 'Customers', COUNT(*) FROM CUSTOMERS  
UNION ALL  
SELECT 'Invoices', COUNT(*) FROM INVOICES  
UNION ALL  
SELECT 'Payments', COUNT(*) FROM PAYMENTS;
```

'CUSTOMERS'	...	COUNT(*)
Customers		50
Invoices		200
Payments		180

-- Checking if there are no matching customer

```
SELECT  
    i.ext_invoice_key, i.ext_customer_key  
FROM  
    Invoices i  
LEFT JOIN  
    Customers c ON i.ext_customer_key = c.ext_customer_key  
WHERE  
    c.ext_customer_key IS NULL;
```

EXT_INVOICE_KEY	EXT_CUSTOMER_KEY
Query produced no results	

-- Payments without a customer

```
SELECT  
    p.ext_payment_key, p.ext_customer_key  
FROM  
    Payments p  
LEFT JOIN  
    Customers c ON p.ext_customer_key = c.ext_customer_key
```

WHERE
c.ext_customer_key IS NULL;

EXT_PAYMENT_KEY	EXT_CUSTOMER_KEY
-----------------	------------------

Query produced no results

-- Invoices with no payments (within 30 days)
SELECT
i.ext_invoice_key, i.ext_customer_key, i.invoice_date, i.amount, i.currency
FROM
Invoices i
LEFT JOIN
Payments p ON i.ext_customer_key = p.ext_customer_key AND ABS(DATEDIFF(day,
i.invoice_date, p.payment_date)) <= 30
WHERE
p.ext_payment_key IS NULL;

EXT_INVOICE_KEY	EXT_CUSTOMER_KEY	...	INVOICE_DATE	AMOUNT	CURRENC
INV0001	CUST039		2024-05-17	4720.26	USD
INV0005	CUST008		2024-01-21	1881.79	USD
INV0006	CUST021		2024-05-21	4861.73	USD
INV0007	CUST039		2024-03-16	4815.99	USD
INV0010	CUST011		2024-06-07	1574.30	USD
INV0019	CUST024		2024-06-18	809.98	USD
INV0025	CUST033		2024-02-29	1264.42	USD
INV0026	CUST012		2024-08-09	3668.26	USD
INV0031	CUST027		2024-06-11	542.42	USD
INV0033	CUST028		2024-04-03	1671.82	USD
INV0034	CUST016		2024-02-11	1013.94	USD
INV0045	CUST004		2024-06-04	4689.98	USD
INV0046	CUST025		2024-06-11	773.85	USD
INV0048	CUST050		2024-05-16	656.02	USD
INV0053	CUST028		2024-04-07	4104.39	USD
INV0055	CUST007		2024-08-24	2695.29	USD
INV0065	CUST004		2024-01-15	4446.72	USD
INV0068	CUST042		2024-01-09	512.29	USD
INV0071	CUST018		2024-04-12	3071.50	USD
INV0077	CUST014		2024-08-12	2788.80	USD
INV0080	CUST015		2024-02-24	1198.92	USD
INV0094	CUST041		2024-05-17	2026.18	USD
INV0098	CUST001		2024-04-19	2562.92	USD
INV0106	CUST024		2024-05-06	967.84	USD
INV0107	CUST011		2024-01-24	4708.25	USD
INV0108	CUST017		2024-07-12	4774.25	USD
INV0109	CUST008		2024-01-13	4582.84	USD
INV0112	CUST033		2024-02-09	4648.76	USD

- List of customers with unpaid invoices.

-- Payments with no invoices (within 30 days)

```
SELECT
  p.ext_payment_key, p.ext_customer_key, p.payment_date, p.amount, p.currency
FROM
  Payments p
LEFT JOIN
  Invoices i ON p.ext_customer_key = i.ext_customer_key AND ABS(DATEDIFF(day,
i.invoice_date, p.payment_date)) <= 30
WHERE
  i.ext_invoice_key IS NULL;
```

EXT_PAYMENT_KEY	EXT_CUSTOMER_KEY	PAYMENT_DATE	AMOUNT	CURRENCY
PAY0079	CUST046	2024-05-05	2408.22	USD
PAY0069	CUST046	2024-05-15	4168.08	USD

- List of payments that doesn't seem to be linked with any invoices.

-- Currency mismatches

```
SELECT
  i.ext_invoice_key, i.currency AS Invoice_Currency, p.ext_payment_key, p.currency AS
Payment_Currency
FROM
  Invoices i
JOIN
  Payments p ON i.ext_customer_key = p.ext_customer_key AND ABS(DATEDIFF(day,
i.invoice_date, p.payment_date)) <= 30
WHERE
  i.currency <> p.currency;
```

EXT_INVOICE_KEY	INVOICE_CURRENCY	EXT_PAYMENT_KEY	PAYMENT_CURRENCY
-----------------	------------------	-----------------	------------------

Query produced no results

- No mismatch in currencies between invoices and payments tables.

-- Overpayment or Underpayment

-- Invoices that have any single payment that doesn't exactly equal the invoice amount

```
SELECT
  i.ext_invoice_key, i.amount AS Invoice_Amount, p.ext_payment_key, p.amount AS
Payment_Amount,
CASE
  WHEN i.amount > p.amount THEN 'Underpaid'
```

```

        WHEN i.amount < p.amount THEN 'Overpaid'
    END AS payment_status
FROM
    Invoices i
JOIN
    Payments p ON i.ext_customer_key = p.ext_customer_key AND ABS(DATEDIFF(day,
i.invoice_date, p.payment_date)) <= 30
WHERE
    p.amount <> i.amount;

```

EXT_INVOICE_KEY	... INVOICE_AMOUNT	EXT_PAYMENT_KEY	PAYMENT_AMOUNT	PAYMENT_STATUS
INV0002	1683.69	PAY0034	1140.90	Underpaid
INV0002	1683.69	PAY0060	1887.68	Overpaid
INV0002	1683.69	PAY0114	1157.52	Underpaid
INV0002	1683.69	PAY0161	1157.52	Underpaid
INV0003	2642.07	PAY0122	4461.03	Overpaid
INV0004	3544.79	PAY0082	2954.78	Underpaid
INV0004	3544.79	PAY0112	3887.11	Overpaid
INV0008	1333.73	PAY0101	1830.10	Overpaid
INV0011	1495.72	PAY0023	4196.30	Overpaid
INV0013	3086.87	PAY0073	3152.16	Overpaid
INV0013	3086.87	PAY0085	3152.16	Overpaid
INV0013	3086.87	PAY0092	3152.16	Overpaid
INV0015	352.25	PAY0022	1901.81	Overpaid
INV0015	352.25	PAY0104	1901.81	Overpaid
INV0015	352.25	PAY0124	1901.81	Overpaid
INV0016	1465.37	PAY0133	181.28	Underpaid
INV0017	4550.50	PAY0078	4399.13	Underpaid
INV0018	1273.85	PAY0081	3134.12	Overpaid
INV0018	1273.85	PAY0172	3393.46	Overpaid
INV0020	2498.32	PAY0119	3420.07	Overpaid
INV0020	2498.32	PAY0167	170.53	Underpaid
INV0021	4929.69	PAY0010	5195.51	Overpaid
INV0021	4929.69	PAY0147	249.45	Underpaid
INV0022	1286.07	PAY0143	957.92	Underpaid
INV0023	3393.46	PAY0075	1273.85	Underpaid
INV0023	3393.46	PAY0081	3134.12	Underpaid
INV0023	3393.46	PAY0116	1273.85	Underpaid
INV0024	3831.94	PAY0070	4324.32	Overpaid

- List of payments which are more or less than the invoice amounts.
- We can also observe that one invoice is linked to multiple payments.

-- Partial Payments

-- Invoices where sum of payments within 30 days equals the invoice amount

SELECT

i.ext_invoice_key, i.amount AS invoice_amount, COUNT(p.ext_payment_key) AS
No_payments, SUM(p.amount) AS total_payment

FROM

Invoices i

JOIN

Payments p ON i.ext_customer_key = p.ext_customer_key AND ABS(DATEDIFF(day,
i.invoice_date, p.payment_date)) BETWEEN 0 AND 30

GROUP BY

i.ext_invoice_key, i.amount

HAVING

SUM(p.amount) = i.amount AND COUNT(p.ext_payment_key) > 1;

EXT_INVOICE_KEY	INVOICE_AMOUNT	NO_PAYMENTS	TOTAL_PAYMENT
Query produced no results			

- No partially settled payments.

-- Canonical tables are created as they represent single source of true data.

-- They hold clean, de-duplicated, and standardized data that is ready for analysis.

-- In our project, the raw CSV data might be messy or may contain duplicate data

-- In order to proceed with better analysis, raw data is merged into canonical tables
(dim_customer, fact_invoice, fact_payment)

-- so that reconciliation rules can be applied consistently.

-- This step ensures Power BI reports are built only on trusted data,

-- avoiding duplicates, errors, and mismatches that exist in raw files.

```
CREATE OR REPLACE TABLE UTRP_SCHEMA.dim_customer (
  customer_id    NUMBER AUTOINCREMENT, -- surrogate key
  ext_customer_key VARCHAR,
  name           VARCHAR,
  email          VARCHAR,
  created_at     TIMESTAMP_NTZ,
  inserted_at    TIMESTAMP_NTZ,
  updated_at     TIMESTAMP_NTZ,
  PRIMARY KEY (customer_id)
);
```

```
CREATE OR REPLACE TABLE UTRP_SCHEMA.fact_invoice (  
  invoice_id    NUMBER AUTOINCREMENT,  
  ext_invoice_key VARCHAR,  
  ext_customer_key VARCHAR,  
  customer_id   NUMBER,  
  invoice_date  DATE,  
  amount        NUMBER(10,2),  
  currency      VARCHAR,  
  status        VARCHAR,  
  source_system VARCHAR,  
  inserted_at   TIMESTAMP_NTZ,  
  updated_at    TIMESTAMP_NTZ,  
  PRIMARY KEY (invoice_id)  
);
```

```
CREATE OR REPLACE TABLE UTRP_SCHEMA.fact_payment (  
  payment_id    NUMBER AUTOINCREMENT,  
  ext_payment_key VARCHAR,  
  ext_customer_key VARCHAR,  
  customer_id   NUMBER,  
  payment_date  DATE,  
  amount        NUMBER(10,2),  
  currency      VARCHAR,  
  payment_method VARCHAR,  
  source_system VARCHAR,  
  inserted_at   TIMESTAMP_NTZ,  
  updated_at    TIMESTAMP_NTZ,  
  PRIMARY KEY (payment_id)  
);
```

```
CREATE OR REPLACE TABLE UTRP_SCHEMA.audit_merge_log (  
  run_id        NUMBER AUTOINCREMENT,  
  object_type   VARCHAR,  
  rows_processed NUMBER,  
  rows_inserted NUMBER,  
  rows_updated  NUMBER,  
  run_by        VARCHAR,  
  run_at        TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),  
  status        VARCHAR,  
  message       VARCHAR  
);
```

- We created canonical tables to make sure we always have clean and reliable data for reconciliation. Instead of building reports directly on raw data which can be messy and duplicated.
- So, we first standardize and store it in these canonical tables so that Power BI dashboards and business users only see the trusted version.

/*

=====

Procedure: sp_merge_invoices

Purpose: Merge (upsert) invoices from staging into canonical fact_invoice table.

Why MERGE?

- Ensures invoices are always clean and up-to-date.
- If invoice is new then insert it.
- If invoice already exists but changed then update it.
- If invoice already exists and unchanged then do nothing.
- This avoids duplicates and stale data.

Why in a stored procedure?

- Automates the full process (count → merge → log → mark as processed).
- Guarantees idempotency (safe to rerun, no duplicates).
- Creates an audit trail (rows processed, inserted, updated).

Why important for this project?

- Canonical tables (fact_invoice, fact_payment, dim_customer) are the trusted single source of truth.
- Reconciliation logic depends on invoices being accurate.
- Finance teams can rely on this clean layer for Power BI dashboards.

=====

*/

```
CREATE OR REPLACE PROCEDURE UTRP_SCHEMA.sp_merge_invoices(run_by STRING)
RETURNS VARIANT
LANGUAGE SQL
EXECUTE AS CALLER
AS
$$
DECLARE
    run_ts      TIMESTAMP_NTZ := CURRENT_TIMESTAMP();
    rows_processed  NUMBER := 0;
    rows_inserted   NUMBER := 0;
    rows_updated    NUMBER := 0;
BEGIN
    -- Count rows to process
    SET rows_processed = (
        SELECT COUNT(*) FROM UTRP_SCHEMA.invoices WHERE processed = FALSE
    );
```



```

-- Merge invoices into canonical fact_invoice
MERGE INTO UTRP_SCHEMA.fact_invoice AS target
USING (
    SELECT s.*, c.customer_id
    FROM UTRP_SCHEMA.invoices s
    LEFT JOIN UTRP_SCHEMA.dim_customer c
        ON s.ext_customer_key = c.ext_customer_key
    WHERE s.processed = FALSE
) AS src
ON target.ext_invoice_key = src.ext_invoice_key
AND target.source_system = src.source_system
WHEN MATCHED AND (
    COALESCE(target.amount,0) <> COALESCE(src.amount,0)
    OR COALESCE(target.status,'') <> COALESCE(src.status,'')
    OR target.invoice_date <> src.invoice_date
)
THEN UPDATE SET
    amount    = src.amount,
    status    = src.status,
    invoice_date = src.invoice_date,
    updated_at = :run_ts
WHEN NOT MATCHED THEN
    INSERT (ext_invoice_key, ext_customer_key, customer_id, invoice_date, amount, currency,
status, source_system, inserted_at, updated_at)
    VALUES (src.ext_invoice_key, src.ext_customer_key, src.customer_id, src.invoice_date,
src.amount, src.currency, src.status, src.source_system, :run_ts, :run_ts);

-- Compute row counts
SET rows_inserted = (
    SELECT COUNT(*) FROM UTRP_SCHEMA.fact_invoice WHERE inserted_at = :run_ts
);
SET rows_updated = (
    SELECT COUNT(*) FROM UTRP_SCHEMA.fact_invoice WHERE updated_at = :run_ts AND
inserted_at <> :run_ts
);

-- Log run
INSERT INTO UTRP_SCHEMA.audit_merge_log(object_type, rows_processed, rows_inserted,
rows_updated, run_by, run_at, status)
VALUES('fact_invoice', :rows_processed, :rows_inserted, :rows_updated, :run_by, :run_ts,
'SUCCESS');

-- Mark invoices as processed

```

```
UPDATE UTRP_SCHEMA.invoices
  SET processed = TRUE
 WHERE processed = FALSE;
```

```
RETURN OBJECT_CONSTRUCT('status','SUCCESS',
                        'processed', :rows_processed,
                        'inserted', :rows_inserted,
                        'updated', :rows_updated);
```

```
EXCEPTION
```

```
  WHEN STATEMENT_ERROR OR EXPRESSION_ERROR THEN
    INSERT INTO UTRP_SCHEMA.audit_merge_log(object_type, rows_processed, rows_inserted,
rows_updated, run_by, run_at, status, message)
      VALUES('fact_invoice', :rows_processed, 0, 0, :run_by, :run_ts, 'FAILED', :sqlerrm);
```

```
    RETURN OBJECT_CONSTRUCT(
      'status','FAILED',
      'error_message', :sqlerrm,
      'sqlstate', :sqlstate
    );
  END;
$$;
```

```
/*
```

```
=====
```

Procedure: sp_merge_customers

Purpose : Merge (upsert) customers from staging into canonical dim_customer table.

Why MERGE?

- Ensures customer records are clean and consistent.
- If customer is new then insert it.
- If customer already exists but details changed then update it.
- If customer already exists and unchanged then do nothing.
- This avoids duplicates and keeps master data reliable.

Why in a stored procedure?

- Automates the end-to-end process (count → merge → log → mark processed).
- Guarantees idempotency (safe to rerun, no duplicates).
- Creates an audit log (rows processed, inserted, updated).

Why important for this project?

- Customer is the key dimension linking invoices and payments.

- Canonical customer table ensures reconciliation always maps invoices and payments to the right entity.
- Provides a single source of truth for Power BI and downstream reporting.

```
=====
*/
CREATE OR REPLACE PROCEDURE UTRP_SCHEMA.sp_merge_customers(run_by STRING)
RETURNS VARIANT
LANGUAGE SQL
EXECUTE AS CALLER
AS
$$
DECLARE
    run_ts      TIMESTAMP_NTZ := CURRENT_TIMESTAMP();
    rows_processed  NUMBER := 0;
    rows_inserted   NUMBER := 0;
    rows_updated    NUMBER := 0;
BEGIN
    -- Count rows to process (assignment context: no colon on left; SQL inside needs : if using var)
    SET rows_processed = (
        SELECT COUNT(*) FROM UTRP_SCHEMA.customers WHERE processed = FALSE
    );

    -- MERGE: note all scripting variables used in SQL are prefixed with :
    MERGE INTO UTRP_SCHEMA.dim_customer AS target
    USING (
        SELECT *
        FROM UTRP_SCHEMA.customers
        WHERE processed = FALSE
    ) AS src
    ON target.ext_customer_key = src.ext_customer_key
    WHEN MATCHED AND (
        COALESCE(target.name, '') <> COALESCE(src.name, '')
        OR COALESCE(target.email, '') <> COALESCE(src.email, '')
    )
    THEN UPDATE SET
        name      = src.name,
        email     = src.email,
        updated_at = :run_ts
    WHEN NOT MATCHED THEN
        INSERT (ext_customer_key, name, email, created_at, inserted_at, updated_at)
        VALUES (src.ext_customer_key, src.name, src.email, src.created_at, :run_ts, :run_ts);

    -- Compute row counts: use :run_ts inside the subqueries (SQL context)
    SET rows_inserted = (
```

```

    SELECT COUNT(*) FROM UTRP_SCHEMA.dim_customer WHERE inserted_at = :run_ts
);
SET rows_updated = (
    SELECT COUNT(*) FROM UTRP_SCHEMA.dim_customer WHERE updated_at = :run_ts AND
inserted_at <> :run_ts
);

-- Log run: use :vars inside the INSERT
INSERT INTO UTRP_SCHEMA.audit_merge_log(
    object_type, rows_processed, rows_inserted, rows_updated, run_by, run_at, status
)
VALUES('dim_customer', :rows_processed, :rows_inserted, :rows_updated, :run_by, :run_ts,
'SUCCESS');

-- Mark staging rows processed
UPDATE UTRP_SCHEMA.customers
    SET processed = TRUE
    WHERE processed = FALSE;

-- Return a JSON-like summary (use :vars)
RETURN OBJECT_CONSTRUCT(
    'status','SUCCESS',
    'processed', :rows_processed,
    'inserted', :rows_inserted,
    'updated', :rows_updated
);

-- Catch-all handler for Snowflake scripting
EXCEPTION
    WHEN STATEMENT_ERROR OR EXPRESSION_ERROR THEN
        INSERT INTO UTRP_SCHEMA.audit_merge_log(
            object_type, rows_processed, rows_inserted, rows_updated, run_by, run_at, status,
message
        )
        VALUES('dim_customer', :rows_processed, 0, 0, :run_by, :run_ts, 'FAILED', :sqlerrm);

    RETURN OBJECT_CONSTRUCT(
        'status','FAILED',
        'error_message', :sqlerrm,
        'sqlstate', :sqlstate
    );
END;
$$;

```

/*

=====

Procedure: sp_merge_payments

**Purpose: Merge (upsert) payments from staging into
canonical fact_payment table.**

Why MERGE?

- Ensures payments are stored once and always up-to-date.
- If payment is new then insert it.
- If payment already exists but details changed then update it.
- If payment already exists and unchanged then do nothing.
- This prevents duplicates and ensures reliable reconciliation.

Why in a stored procedure?

- Automates the process end-to-end (count → merge → log → mark processed).
- Idempotent design means you can rerun safely without duplication.
- Audit logging ensures every run is tracked for compliance.

Why important for this project?

- Payments are the core fact table to reconcile against invoices.
- Clean, canonical payments data allows accurate matching rules (exact match, partial, over/under, currency mismatch).
- Provides the foundation for exception reporting in Power BI.

=====

*/

CREATE OR REPLACE PROCEDURE UTRP_SCHEMA.sp_merge_payments(run_by STRING)

RETURNS VARIANT

LANGUAGE SQL

EXECUTE AS CALLER

AS

\$\$

DECLARE

run_ts TIMESTAMP_NTZ := CURRENT_TIMESTAMP();

rows_processed NUMBER := 0;

rows_inserted NUMBER := 0;

rows_updated NUMBER := 0;

BEGIN

-- Count rows to process

SET rows_processed = (

 SELECT COUNT(*) FROM UTRP_SCHEMA.payments WHERE processed = FALSE

);

-- Merge payments into canonical fact_payment

```

MERGE INTO UTRP_SCHEMA.fact_payment AS target
USING (
    SELECT s.*, c.customer_id
    FROM UTRP_SCHEMA.payments s
    LEFT JOIN UTRP_SCHEMA.dim_customer c
        ON s.ext_customer_key = c.ext_customer_key
    WHERE s.processed = FALSE
) AS src
ON target.ext_payment_key = src.ext_payment_key
AND target.source_system = src.source_system
WHEN MATCHED AND (
    COALESCE(target.amount,0) <> COALESCE(src.amount,0)
    OR COALESCE(target.payment_method,'') <> COALESCE(src.payment_method,'')
    OR target.payment_date <> src.payment_date
)
THEN UPDATE SET
    amount      = src.amount,
    payment_method = src.payment_method,
    payment_date = src.payment_date,
    updated_at   = :run_ts
WHEN NOT MATCHED THEN
    INSERT (ext_payment_key, ext_customer_key, customer_id, payment_date, amount,
currency, payment_method, source_system, inserted_at, updated_at)
    VALUES (src.ext_payment_key, src.ext_customer_key, src.customer_id, src.payment_date,
src.amount, src.currency, src.payment_method, src.source_system, :run_ts, :run_ts);

-- Compute row counts
SET rows_inserted = (
    SELECT COUNT(*) FROM UTRP_SCHEMA.fact_payment WHERE inserted_at = :run_ts
);
SET rows_updated = (
    SELECT COUNT(*) FROM UTRP_SCHEMA.fact_payment WHERE updated_at = :run_ts AND
inserted_at <> :run_ts
);

-- Log run
INSERT INTO UTRP_SCHEMA.audit_merge_log(object_type, rows_processed, rows_inserted,
rows_updated, run_by, run_at, status)
VALUES('fact_payment', :rows_processed, :rows_inserted, :rows_updated, :run_by, :run_ts,
'SUCCESS');

-- Mark payments as processed
UPDATE UTRP_SCHEMA.payments
SET processed = TRUE

```

```

WHERE processed = FALSE;

RETURN OBJECT_CONSTRUCT('status','SUCCESS',
                        'processed',:rows_processed,
                        'inserted',:rows_inserted,
                        'updated',:rows_updated);

EXCEPTION
  WHEN STATEMENT_ERROR OR EXPRESSION_ERROR THEN
    INSERT INTO UTRP_SCHEMA.audit_merge_log(object_type, rows_processed, rows_inserted,
rows_updated, run_by, run_at, status, message)
    VALUES('fact_payment',:rows_processed, 0, 0, :run_by, :run_ts, 'FAILED', :sqlerrm);

RETURN OBJECT_CONSTRUCT(
  'status','FAILED',
  'error_message',:sqlerrm,
  'sqlstate',:sqlstate
);
END;
$$;

```

What does the query actually do?

- Count rows to process - Check how many new/unprocessed invoices are waiting in the staging table.
- Match staging to canonical:
 - If invoice already exists but with changes then update it.
 - If invoice doesn't exist then insert it.
 - If invoice exists and nothing changed then leave it.
- Log results - Insert a row into audit_merge_log with how many rows were processed, inserted, or updated.
- Mark staging rows as processed - So they won't be picked up again in the next run.
- Return a summary - Success/failure + row counts.

RECONCILIATION ENGINE

-- Reconciliation Engine Design

```

CREATE OR REPLACE TABLE UTRP_SCHEMA.recon_invoice_payment (
  recon_id    STRING DEFAULT UUID_STRING(),
  ext_invoice_key STRING,
  customer_id  STRING,
  invoice_amount  NUMBER,
  payment_total  NUMBER,
  difference    NUMBER,

```

```

    match_status STRING, -- e.g. 'MATCHED', 'PARTIAL', 'OVERPAID', 'UNDERPAID', 'UNPAID'
    run_at        TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP()
);

```

```

CREATE OR REPLACE TABLE UTRP_SCHEMA.recon_exceptions (
    exception_id STRING DEFAULT UUID_STRING(),
    ext_invoice_key STRING,
    customer_id  STRING,
    issue_type   STRING, -- e.g. 'NO_MATCH', 'PARTIAL', 'OVERPAYMENT',
'MISSING_CUSTOMER'
    issue_details STRING,
    suggested_action STRING,
    resolved      BOOLEAN DEFAULT FALSE,
    run_at        TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP()
);

```

-- Reconciliation Engine

```

CREATE OR REPLACE PROCEDURE UTRP_SCHEMA.sp_reconcile_invoices_payments(run_by
STRING)
RETURNS VARIANT
LANGUAGE SQL
EXECUTE AS CALLER
AS
$$
DECLARE
    run_id STRING DEFAULT UUID_STRING(); -- unique identifier for this run
    run_ts TIMESTAMP_NTZ := CURRENT_TIMESTAMP();
    rows_processed NUMBER := 0;
    rows_matched NUMBER := 0;
    rows_exceptions NUMBER := 0;
BEGIN
    -- Insert reconciliation results
    INSERT INTO UTRP_SCHEMA.recon_invoice_payment
    (
        ext_invoice_key, customer_id, invoice_amount, payment_total, difference, match_status,
run_at, run_id
    )
    SELECT
        fi.ext_invoice_key,
        fi.customer_id,
        fi.amount AS invoice_amount,
        COALESCE(SUM(fp.amount),0) AS payment_total,
        fi.amount - COALESCE(SUM(fp.amount),0) AS difference,
        CASE

```



```

        WHEN COALESCE(SUM(fp.amount),0) = fi.amount THEN 'MATCHED'
        WHEN COALESCE(SUM(fp.amount),0) = 0 THEN 'UNPAID'
        WHEN COALESCE(SUM(fp.amount),0) < fi.amount THEN 'PARTIAL'
        WHEN COALESCE(SUM(fp.amount),0) > fi.amount THEN 'OVERPAID'
        ELSE 'UNKNOWN'
    END AS match_status,
    :run_ts,
    :run_id
FROM UTRP_SCHEMA.fact_invoice fi
LEFT JOIN UTRP_SCHEMA.fact_payment fp
    ON fi.customer_id = fp.customer_id
GROUP BY fi.ext_invoice_key, fi.customer_id, fi.amount;

-- Count reconciled rows for this run
SELECT COUNT(*) INTO :rows_processed
FROM UTRP_SCHEMA.recon_invoice_payment
WHERE run_id = :run_id;

SELECT COUNT(*) INTO :rows_matched
FROM UTRP_SCHEMA.recon_invoice_payment
WHERE run_id = :run_id
AND match_status = 'MATCHED';

-- Insert exceptions for non-matches
INSERT INTO UTRP_SCHEMA.recon_exceptions
(
    ext_invoice_key, customer_id, issue_type, issue_details, suggested_action, run_at, run_id
)
SELECT
    r.ext_invoice_key,
    r.customer_id,
    CASE r.match_status
        WHEN 'UNPAID' THEN 'NO_PAYMENT'
        WHEN 'PARTIAL' THEN 'PARTIAL_PAYMENT'
        WHEN 'OVERPAID' THEN 'OVERPAYMENT'
        ELSE 'UNKNOWN'
    END AS issue_type,
    'Invoice amount = ' || r.invoice_amount || ', Payments = ' || r.payment_total || ',
Difference = ' || r.difference AS issue_details,
    CASE r.match_status
        WHEN 'UNPAID' THEN 'Investigate missing payment or follow up with customer'
        WHEN 'PARTIAL' THEN 'Review outstanding balance and initiate collection'
        WHEN 'OVERPAID' THEN 'Check for duplicate/refund required'
        ELSE 'Manual review needed'
    
```

```

        END AS suggested_action,
        :run_ts,
        :run_id
FROM UTRP_SCHEMA.recon_invoice_payment r
WHERE r.run_id = :run_id
      AND r.match_status <> 'MATCHED';

-- Count exceptions
SELECT COUNT(*) INTO :rows_exceptions
FROM UTRP_SCHEMA.recon_exceptions
WHERE run_id = :run_id;

-- Return proper JSON summary
RETURN OBJECT_CONSTRUCT(
    'status','SUCCESS',
    'processed', :rows_processed,
    'matched', :rows_matched,
    'exceptions', :rows_exceptions,
    'run_id', :run_id,
    'run_by', :run_by,
    'run_at', :run_ts
);

EXCEPTION
WHEN STATEMENT_ERROR OR EXPRESSION_ERROR THEN
RETURN OBJECT_CONSTRUCT(
    'status','FAILED',
    'error_message', :sqlerrm,
    'sqlstate', :sqlstate,
    'run_by', :run_by,
    'run_at', :run_ts,
    'run_id', :run_id
);
END;
$$;

-- Calling Stored Procedure
CALL UTRP_SCHEMA.sp_reconcile_invoices_payments('shijin');

-- Adding columns in both the recon tables
ALTER TABLE UTRP_SCHEMA.recon_invoice_payment
ADD COLUMN run_id STRING;

```

```
ALTER TABLE UTRP_SCHEMA.recon_exceptions  
ADD COLUMN run_id STRING;
```

```
ALTER TABLE UTRP_SCHEMA.fact_payment  
ADD COLUMN invoice_id INT;
```

Summary

- In most real-world finance and IT systems, invoices and payments don't always line up perfectly. Some invoices are fully paid, some are unpaid, some are partially paid, and sometimes customers even overpay.
- Manually tracking these mismatches is time-consuming and error-prone for finance teams.
- To solve this, we built a Reconciliation Engine that automatically compares invoices against payments and classifies the results.

How it works

1. Join Invoices and Payments

- We linked invoices from fact_invoice with payments from fact_payment based on customer_id.
- For each invoice, we calculated the total payments made by that customer.

2. Calculate Differences

- We compared invoice amount vs total payments.
- The difference tells us whether the invoice is fully settled or has issues.

3. Classify Match Status

- If payments = invoice amount → MATCHED
- If no payments → UNPAID
- If payments < invoice amount → PARTIAL
- If payments > invoice amount → OVERPAID

4. Store Results

- All reconciliation results are stored in recon_invoice_payment.
- Any problematic cases (UNPAID, PARTIAL, OVERPAID) are logged into recon_exceptions with:
- Issue Type (e.g., No Payment, Partial, Overpayment)

- Issue Details (amounts and differences)
- Suggested Action (e.g., “Follow up with customer”, “Check for duplicate/refund”)

5. Run Tracking with run_id

- Each reconciliation run is tagged with a unique run_id.
- This makes it easy to audit and analyze multiple runs.

Why these matters in the industry?

- Finance teams don’t need to manually compare invoice and payment records.
- The engine automatically highlights mismatches and gives clear next steps.
- Exceptions are captured in a structured way, ready for reporting in Power BI.
- Improves accuracy, saves time, and reduces financial risk.

With this, we automated the invoice vs payment matching process. Instead of relying on manual checks, our reconciliation engine now classifies every invoice into Matched, Unpaid, Partial, or Overpaid, and logs exceptions for review.

Reconciliation Engine

